

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

CRC magic tricks

dram

2021-12-29

About me

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- @dramforever on most random platforms
 - GitHub, Twitter, ...
- <https://dram.page>
- Call me 'dram'

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Background

Galois field GF(2)

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Finite field
- Elements: 0, 1
- Addition is XOR
- Multiplication is AND

Notable properties:

- $2 = 0$
- $a + b = a - b$

Polynomials in GF(2)

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

$$p(x) = \sum_{n=0\dots d} a_n x^n$$

- $a_n \in \text{GF}(2)$ are the *coefficients*
- $\deg(p(x))$: Power of highest power term with non-zero coefficient
- x is just a symbol
 - Polynomials are not functions

Polynomial addition

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

– Addition: $(x^3 + x) + (x + 1)$

$$\begin{array}{rccccccc} & x^3 & & + & x & & \\ & & & + & x & + & 1 \\ \hline = & x^3 & & & & + & 1 \end{array}$$

Polynomial multiplication

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

– Multiplication: $(x + 1)(x^3 + x + 1)$

$$\begin{array}{r} x^4 + x^3 \\ x^2 + x \\ x + 1 \\ \hline = x^4 + x^3 + x^2 + 1 \end{array}$$

Polynomial Euclidean division

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Given $a(x)$ and $b(x)$, there is unique $(q(x), r(x))$ such that $\deg(r(x)) < \deg(b(x))$ and:

$$a(x) = b(x) \cdot q(x) + r(x)$$

Some shorthands:

- Quotient: $q(x) = a(x) \operatorname{div} b(x)$
 - (Note: Gopal et al. (2009) writes this as $\lfloor a(x)/b(x) \rfloor$)
- Remainder: $r(x) = a(x) \bmod b(x)$

Polynomial GCD

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- $b(x)$ *divides* $a(x)$ iff $a(x) \bmod b(x) = 0$
- $\gcd(a(x), b(x))$ is the unique largest-degree polynomial $g(x)$ that divides both $a(x)$ and $b(x)$

Euclidean algorithm works for polynomial GCD

```
def poly_gcd(a, b):  
    while b != 0:  
        a, b = b, poly_mod(a, b)  
    return a
```

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Cyclic redundancy check

Cyclic redundancy check

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- An family error *detecting* codes
- Based on polynomials in $\text{GF}(2)$
- Not cryptographically secure *at all*
- Commonly called CRC- N for a CRC with an N -bit check sequence
- No single standard, parameters vary greatly
 - (For a catalogue of various CRCs see Cook (2021a))

CRC implemented as LFSRs

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Input message as bit-stream
 - For each byte, put LSB first¹
 - "\xf0\x30" → 0000 1111 0000 1100

```
def crc(message):  
    crc = INIT  
  
    for b in message:  
        crc ^= b  
        if crc & 1: crc = (crc >> 1) ^ TAP  
        else:      crc = crc >> 1  
  
    return crc ^ FINAL
```

¹Some implementations use other bit orders.

Bit-streams and polynomials

CRC tricks

dram

Background

Cyclic
redundancy
check


Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Least significant bit = First transmitted bit = Highest power term
- Parameters INIT, TAP, FINAL are *bit-streams*
- The register crc is a *bit-stream*



Back to the CRC code

- Suppose we're working with an N -bit CRC

```
def crc(message):  
    crc = INIT  
  
    for b in message:  
        #  $crc \leftarrow crc + b x^{(N - 1)}$   
        crc ^= b  
  
        # if  $crc$  has  $x^{(N - 1)}$  term  
        if crc & 1:  
            #  $crc \leftarrow crc x + x^N + TAP$   
            # (Right shift discards  $x^{(N - 1)}$  term instead of turning it into  $x^N$ )  
            crc = (crc >> 1) ^ TAP  
        else:  
            #  $crc \leftarrow crc x$   
            crc = crc >> 1  
  
    return crc ^ FINAL
```

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

CRC computation, but polynomials

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

The main loop:

- For each bit b in message:
 - $\text{crc} \leftarrow \text{crc} + bx^{N-1}$
 - If the x^{N-1} coefficient of crc is 1, then $\text{crc} \leftarrow \text{crc} \cdot x + x^N + \text{tap}$
 - Else: $\text{crc} \leftarrow \text{crc} \cdot x$

Simplified:

- For each bit b in message:
 - $\text{crc} \leftarrow \text{crc} \cdot x + bx^N$
 - If the x^N coefficient of crc is 1,
 - Then $\text{crc} \leftarrow \text{crc} + x^N + \text{tap}$

Why is the x^N here

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Example: `crc = 0x3`, LSB of `crc` is 1
- $\text{crc} = x^{N-2} + x^{N-1}$

After shifting

- $\text{crc} \gg 3 = 0x1$
- $\text{crc} \cdot x = x^{N-1} + x^N$
- $\text{crc} \cdot x + x^N = x^{N-1}$

What does the algorithm do?

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Without TAP:

- $\text{crc}^* \leftarrow \text{init}$
- For each bit b in message:
 - $\text{crc}^* \leftarrow \text{crc}^* \cdot x + bx^N$
 - ~~If the x^N coefficient of crc^* is 1,~~
 - ~~Then $\text{crc}^* \leftarrow \text{crc}^* + x^N + \text{tap}$~~
- Return $\text{crc}^* + \text{final}$

Let m be the message, with length in bits L , then the result is:

$$\text{crc}^* = mx^N + \text{init} \cdot x^L + \text{final}$$

What does the algorithm do?

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Claim: $\text{crc} \equiv \text{crc}^* \pmod{(\text{tap} + x^N)}$
- $\deg(\text{crc}) < N$

Therefore:

$$\begin{aligned}\text{crc} &= \text{crc}^* \bmod (\text{tap} + x^N) \\ &= (mx^N + \text{init} \cdot x^L + \text{final}) \bmod (\text{tap} + x^N)\end{aligned}$$

Shorter symbols

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

From now on, we'll use these symbols consistently:

- L (length) is the message length in bits
- m (message) is the message bit string as a polynomial
- N is the length of the CRC
- r (remainder) is crc
- F is final
- I is init
- P (polynomial) is tap + x^N

$$r = (mx^N + Ix^L + F) \bmod P$$

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Efficient CRC with CLMUL

Carryless multiplication

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- ‘Carryless multiplication’ operation
 - Intel: `pclmulq`
 - ARM: `pmull`
 - RISC-V (Zbc): `clmul{,h,r}`
- Much faster than software loop
- Also much faster than div/mod
- Directly corresponds to register-sized $GF(2)$ polynomial multiplication

Efficient CRC by folding

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Method described in Intel whitepaper (Gopal et al., 2009)
 - *Fast CRC Computation for Generic Polynomials Using PCLMULQDQ Instruction*
- Keep intermediate result a of $2N$ bits, $a \equiv mx^N + Ix^L \pmod{P}$
- Read message in N -bit chunks, updating a if needed
- Calculate $(a \bmod P) + F$ for final result
- Using precomputed constants, avoids (dynamic) polynomial div/mod entirely.

Folding step

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Read message in N -bit chunks m_k (so $\deg(m_k) < N$), new a should be:

$$\begin{aligned} a^* &\equiv (mx^N) \cdot x^N + m_k x^N + Ix^{L+N} \pmod{P} \\ &\equiv (a + m_k)x^N \pmod{P} \end{aligned}$$

- At each iteration we need to ensure $\deg(a^*) < 2N$
- Split a into ‘high N terms’ and ‘low N terms’, $a = a_H x^N + a_L$, $\deg(a_L) < N$
- $a^* = (a_L + m_k)x^N + a_H(x^{2N} \bmod P)$
- $(x^{2N} \bmod P)$ can be precomputed

Barret reduction

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- We need to find $a \bmod P$
- Suppose $a = Pq + r$, where $\deg(q), \deg(r) < N$
- Let $\mu = x^{2N} \operatorname{div} P$, then $\deg(\mu) = N$
- Property: $\deg(x^{2N} + \mu P) < N$

Barret reduction:

- Let $t = (a \operatorname{div} x^N) \cdot \mu$, then

$$\begin{aligned} t &= (a \operatorname{div} x^N) \cdot \mu \\ &= (\mu Pq \operatorname{div} x^N) + (\mu r \operatorname{div} x^N) \\ &= (((x^{2N} + o(x^N)) \cdot q) \operatorname{div} x^N) + o(x^N) \\ &= x^N \cdot q + o(x^N) \end{aligned}$$

- Therefore $q = t \operatorname{div} x^N$

Barret reduction

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- $r = a + Pq$
- $r + F$ is our final CRC
- μ is precomputed
- (Note: $u \operatorname{div} x^N$ is not really a division, just takes ‘higher half’)

Slightly simplifying μ

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- (Used by Wolf (2019), also in historic RISC-V Bitmanip spec (Wolf, 2021))
- μ is degree N , which does not fit in N bits
- Use $\mu \operatorname{div} x$ instead of μ . Let c_0 be the constant term of μ
- $x^{2N} + x(\mu \operatorname{div} x)P = o(x^N) + c_0P = o(x^{N+1})$
- $x^{2N-1} + (\mu \operatorname{div} x)P = o(x^N)$

$$\begin{aligned}t_1 &= (a \operatorname{div} x^N) \cdot (\mu \operatorname{div} x) \\&= (((x^{2N-1} + o(x^N)) \cdot q) \operatorname{div} x^N) + o(x^{N-1}) \\&= x^{N-1}q + o(x^{N-1})\end{aligned}$$

- Still works: $q = t_1 \operatorname{div} x^{N-1}$
- (Noted by Kutenin (2021) that some implementations of the same CRC-32 differ in the constant term of μ)

Bit reversed CLMUL

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Store N -bit result in N -bit GPR
- Store $2N$ -bit intermediate in two N -bit GPRs
- Polynomial in register is bit reversed
 - LSB = Highest power term

How do we calculate bit-reversed CLMUL?

- Two N -bit inputs, $(2N - 1)$ -bit result
- CLMUL is symmetric:

$$\text{rev}_{2N-1}(\text{clmul}(a, b)) = \text{clmul}(\text{rev}_N(a), \text{rev}_N(b))$$

CLMUL is symmetric

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Highest power term?

$$\begin{array}{rcccccccc} & x^4 & + & x^3 & & & & & \\ & & & & & x^2 & + & x^1 & \\ & & & & & & & x^1 & + & x^0 \\ \hline = & x^4 & + & x^3 & + & x^2 & & & + & x^0 \end{array}$$

- Lowest power term?

$$\begin{array}{rcccccccc} & x^0 & + & x^1 & & & & & \\ & & & & & x^2 & + & x^3 & \\ & & & & & & & x^3 & + & x^4 \\ \hline = & x^0 & + & x^1 & + & x^2 & & & + & x^4 \end{array}$$

CLMUL instructions

CRC tricks

dram

Background

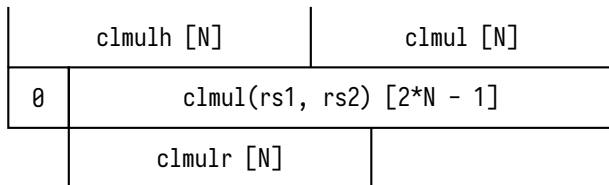
Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion



- In bit reversed representation:
 - $a \cdot b$, low half: clmul rd, rs1, rs2, high half: clmul ; slli 1
 - $(a \cdot b) \text{ div } x^N$ is clmul ; slli 1
 - $(a \cdot b) \text{ div } x^{N-1}$ is clmul

Further speedups

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- pclmulqdq can help handle 128-bit chunks
- Handling multiple (e.g. 4) chunks in parallel
 - Modern processors have many CLMUL units, to keep up with AES for GCM

All these are described in the Intel whitepaper (Gopal et al., 2009).

Table-based CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Without CLMUL, make a table of $T(k) = (k \cdot x^N) \bmod P$ for all $\deg(k) < C$
- Chunk size e.g. $C = 8$, each octet as a chunk
- Let $r = r_H x^{N-T} + r_L$, $\deg(r_L) < N - T$

$$\begin{aligned} r^* &= ((r_H + m_k) \cdot x^N) \bmod P + r_L \\ &= T(r_H + m_k) + r_L \end{aligned}$$

- C is usually very small, as table requires $N \cdot 2^C$ bits

Summary of computing CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- LFSR
- Bit-reversed representation
- Speeding up CRC with CLMUL
 - Folding
 - Barret reduction
- Using RVZbc

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Preimage attacks

'Security' of CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- *CRC has no cryptographic security at all*
- CRC only intends to protect against inadvertent changes, especially those occurring during transmission and storage

Linear/affine properties of CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- For two messages with identical length:

$$r_1 = (m_1x^N + Ix^L + F) \bmod P$$

$$r_2 = (m_2x^N + Ix^L + F) \bmod P$$

- Adding the two gives:

$$r_1 + r_2 = (m_1 + m_2)x^N \bmod P$$

- CRCs are *affine*

Combining messages with XOR

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If $m_1 \oplus m_2 = m_3 \oplus m_4$
- Then $\text{crc}(m_1) \oplus \text{crc}(m_2) = \text{crc}(m_3) \oplus \text{crc}(m_4)$
- In particular, let z be the all-zeros messages with same length as m_k
- Then $\text{crc}(m_1 \oplus m_2) = \text{crc}(m_1) \oplus \text{crc}(m_2) \oplus \text{crc}(z)$

Finding constrained preimages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Example: Find integer i such that its ASCII representation hashes to (known) r :

`crc32(str(i)) = r`

- Application: (Does anyone know what this is?)

`<d p="507.77900,1,25,16777215,1640704211,0,4e291766,59565938560754176,10">HDMI orz</d>`

- Find uid (up to around 10^9) such that:

`crc(str(uid)) = 0x4e291766`

Finding constrained preimages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Simple case, fixed length 9
- Meet-in-the-middle (Dot means $\backslash 0$):

```
crc("123456789")  
=  
  crc("12345....")  
  ^ crc(".....6789")  
  ^ crc(".....")
```

- Generate all 10^5 possible 'high parts', all 10^4 possible 'low parts'
- Hash table of 'high parts' `high_table[0x770a59bd] = "12345...."`
- Array of 'low parts' `low_table[i] = (".....6789", 0x5af77435)`
- `crc(".....") = 0xe60914ae`

Finding constrained preimages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

```
from zlib import crc32

low_table = [
    (i, crc32(b'\x00' * 5 + str(i).encode()))
    for i in range(10**4)
]

high_table = {
    crc32(str(i).encode() + b'\x00' * 4) : i
    for i in range(10**5)
}

def find_num(target):
    crc_z = crc32(b'\x00' * 9)
    for low_num, low_hash in low_table:
        expect = low_hash ^ target ^ crc_z
        if expect in high_table:
            high_num = high_table[expect]
            return high_num * (10**4) + low_num
```

Preimage exercise

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- (Python: `zlib.crc32`)
- Find integer i such that `zlib.crc32(str(i)) == 0x4e291766`
- $0 < i < 10^{**9}$

Bit-flipping messages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Several ‘free positions’ in a message where we can change the bits to anything
- Pick values for these bits such that the message has desired CRC

Use cases:

- Modifying a file so it hashes to interesting values
- Tamper with file without CRC changing

Bit-flipping messages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- For each ‘free position’
- Flipping bit $k \leftrightarrow$ CRC gets bitwise-xor by a_k
- We have N positions we can flip a bit
- Let $d_k = 1$ if flip position k , 0 if no flip

$$\text{orig} \oplus \text{target} = \bigoplus_{k=0 \dots N-1} d_k \cdot a_k$$

- Find linear combination of a_k that gives $\text{orig} \oplus \text{target}$
- Solvable with Gaussian elimination

Bit-flipping messages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Notable features
 - If all we need is where to flip bits, the original message is *not required*
- In general, for CRC- N , N free positions are needed
- We're going to use this GF(2) polynomial linear equation solver later

Summary of preimage attacks

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Constrained charset: Meet-in-the-middle
- Bit flips: Gaussian elimination

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Reverse engineering

CRC without unknown parameters

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Worse case: we only know it's a CRC with width N
- We have a few known (m_k, r_k) pairs
- Rocksoft Model Algorithm parameters (Williams, 1993)
 - N is width
 - P is poly
 - I is init
 - F is xorout
 - refin and refout
 - Bit order of input/output
 - Both assumed to be true here
 - Otherwise, only takes 4 tries to test all combinations
 - name and check are irrelevant
- Used in the CRC catalogue (Cook, 2021a)

How do we even know it's a CRC?

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Use the affine property
- If $m_1 \oplus m_2 = m_3 \oplus m_4$
- Then $\text{crc}(m_1) \oplus \text{crc}(m_2) = \text{crc}(m_3) \oplus \text{crc}(m_4)$

Check for all pairs you can find

- Consecutive numbers or related strings often have linear relationships

Can also help uncover incorrect (m, r) pairs

Messages with CRC appended

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Most often, the CRC comes right after the message
- In terms of polynomials:
 - Message: m
 - Append N zeros: mx^N
 - Set to CRC value: $mx^N + r$

Known message and CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

$$r \equiv mx^N + Ix^L + F \pmod{P}$$

- Known m and r :
- Moving knowns to one side:

$$mx^N + r \equiv Ix^L + F \pmod{P}$$

Comparing equal length messages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If we have two known $m_k x^N + r_k$ with m_1 and m_2 having equal length L

$$m_1 x^N + r_1 \equiv Ix^L + F \pmod{P}$$

$$m_2 x^N + r_2 \equiv Ix^L + F \pmod{P}$$

- Adding the two gives:

$$(m_1 x^N + r_1) + (m_2 x^N + r_2) \equiv 0 \pmod{P}$$

- Shorthand: $v_k = m_k x^N + r_k$

'Difference messages'

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If $L_j = L_k$ then P divides $v_j + v_k$
- Finding all linear independent pairs of (j, k) :
 - Sort all messages by length
 - Find all adjacent equal-length pairs

We have several $v_j + v_k$ that are 'multiples' of P

- Next task: Find degree N polynomial P
- ... such that P divides all $v_j + v_k$ where $L_j = L_k$

CRC RevEng

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- CRC RevEng (Cook, 2021b) can find parameters for CRCs
- Algorithm for finding P : Brute force
 - Search through all polynomials with degree N , and with constant-term
 - Check if divides all differences
- Optimization:
 - If a certain $\deg(v_j + v_k) < 2N$, search for the smaller factor of $v_j + v_k$ instead

An alternative...?

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- P divides all $v_j + v_k$ where $L_j = L_k$
- $\Leftrightarrow P$ divides the GCD of all such $v_j + v_k$

Taking a few GCDs may quickly isolate P :

- The GCD turns out to have degree N , then it is P
- The GCD has a degree slightly larger than N , use ‘find smaller factor’ method
- The GCD still has high degree... Try with more samples

(Even if we don't know N , we can guess that it's the degree of the GCD.)

An example

CRC tricks

dram

- Real data captured from a bus on an adjustable desk:

m1, r1 = AAFF 0040 2EEC

m2, r2 = AAFF 0060 2964

m3, r3 = AAFF 0050 2B08

$$(m_1x^{16} + r_1) + (m_2x^{16} + r_2) = x^{18} + x^{15} + x^{14} + x^{13} + x^4 + x^0$$

$$(m_1x^{16} + r_1) + (m_3x^{16} + r_3) = x^{19} + x^{15} + x^{13} + x^5 + x^2 + x^1 + x^0$$

- Calculating the GCD gives:

$$\gcd(\dots, \dots) = x^{16} + x^{14} + x^{13} + x^2 + x^0$$

- Conclusion: Probably CRC-16 with this P

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Comparing differing-length messages

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If we have two known $m_k x^N + r_k$ with m_1 and m_2 having different lengths L_1 and L_2 respectively

$$m_1 x^N + r_1 \equiv I x^{L_1} + F \pmod{P}$$

$$m_2 x^N + r_2 \equiv I x^{L_2} + F \pmod{P}$$

- Adding the two gives:

$$(m_1 x^N + r_1) + (m_2 x^N + r_2) \equiv I(x^{L_1} + x^{L_2}) \pmod{P}$$

Solving for initial value

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

$$I(x^{L_1} + x^{L_2}) \bmod P = ((m_1 x^N + r_1) + (m_2 x^N + r_2)) \bmod P$$

– Suppose:

$$I = \sum_{k=0 \dots N-1} a_k x^k$$

Then:

$$\begin{aligned} & \sum_{k=0 \dots N-1} a_k (x^k (x^{L_1} + x^{L_2}) \bmod P) \\ &= ((m_1 x^N + r_1) + (m_2 x^N + r_2)) \bmod P \end{aligned}$$

Solving for init

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Finding a linear combination of $(x^k(x^{L_1} + x^{L_2})) \bmod P$ summing to $((m_1x^N + r_1) + (m_2x^N + r_2)) \bmod P$
- It's Gaussian elimination again

Example:

m1, r1 = AAFF 0040 2EEC

m2, r2 = AAFF 040E02 0450

- $L_1 = 32, L_2 = 40$
- Solving gives `init = 0xffff`

Non-unique init values

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- There might not be a unique solutions for I , as noted by Ewing (2010)
- Some CRC polynomials may have multiple equivalent (I, F) pairs
- These polynomials are reducible, i.e. have non-trivial factors

Firstly, two notable properties of $x + 1$

- $x^M + 1 = (x + 1)(x^{M-1} + x^{M-2} + \dots + x + 1)$
- If M is a power of 2, then $x^M + 1 = (x + 1)^M$

Non-unique init values

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If I_1 and I_2 for, then for *all* pairs of natural numbers (L_1, L_2) ,

$$I_1(x^{L_1} + x^{L_2}) \equiv I_2(x^{L_1} + x^{L_2}) \pmod{P}$$

Or equivalently:

$$(I_1 + I_2)(x^{L_1} + x^{L_2}) \equiv 0 \pmod{P}$$

Let $I^* = I_1 + I_2$

Conditions for non-unique init values

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- For *all* pairs of natural numbers (L_1, L_2) , without loss of generality assuming $L_1 > L_2$,

$$I^*(x^{L_1} + x^{L_2}) \equiv 0 \pmod{P}$$

- Bezout's theorem: for all a and b , there exists u and v such that

$$au + bv = \gcd(a, b)$$

- If $I^*a \equiv 0 \pmod{P}$ and $I^*b \equiv 0 \pmod{P}$
- Then $I^*(au + bv) \equiv 0 \pmod{P}$
- Therefore $I^*\gcd(a, b) \equiv 0 \pmod{P}$

What is the GCD then?

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- We need to find I^* such that

$$I^*g \equiv 0 \pmod{P}$$

- Where g is the GCD of all polynomials of the form $x^{L_1} + x^{L_2}$

$$\begin{aligned}x^{L_1} + x^{L_2} &= (x^{L_1-L_2} + 1)x^{L_2} \\ &= (x + 1)(x^{L_1-L_2-1} + \dots + 1)x^{L_2}\end{aligned}$$

- If $L_1 - L_2 = 1$ then $x^{L_1-L_2-1} + \dots + 1 = 1$
- If $L_2 = 0$ then $x^{L_2} = 1$
- $g = x + 1$

If lengths are constrained

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If all lengths L are multiples of some power-of-two ‘byte width’, say $Z = 2^w$
 - $Z = 8$ for octets
 - $Z = 1$ for bit streams

$$\begin{aligned}x^{L_1} + x^{L_2} &= (x^{L_1-L_2} + 1)x^{L_2} \\&= (x^{(L_1-L_2)/Z} + 1)^Z x^{L_2} \\&= (x + 1)^Z (x^{(L_1-L_2)/Z-1} + \dots + 1)^Z x^{L_2}\end{aligned}$$

- Similarly, we have $g = (x + 1)^Z$

Reducible poly values

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- If P has a factor $(x + 1)^f$, then $P/(x + 1)^{\min\{f, Z\}}$ is a valid I^*
- Given valid I , any other $I + uI^*$ is also a valid initial value

Examples of reducible poly

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

ModBus CRC-16, $P = x^{16} + x^{15} + x^2 + 1$

- $I^* = x^{15} + x^1 + x^0$ is a multiple of $(x + 1)$
- Valid init values: 0xffff (standard), 0x3ffe
- Difference is: 0xc001, same as noted by Ewing (2010)

Examples of reducible poly

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Go crc64 package², ECMA polynomial

- (Wrong bit order for ECMA 182 CRC-64)
- Polynomial is multiple of $(x + 1)^2$
- Valid init values: 0xffffffffffffffff (standard), 0x0b8fb9ee4606a6fd, 0x71b79ae69afa0a7c, 0x85c7dcf72303537e

In general, there are $2^{\min\{f, Z\}}$ valid (I, F) pairs, because smallest I^* has degree $N - \min\{f, Z\}$, so $\deg u < \min\{f, Z\}$

²<https://pkg.go.dev/hash/crc64>

Solving for xorout

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Going back to any (m, r)

$$r \equiv mx^N + Ix^L + F \pmod{P}$$

- Solving for F is pretty easy now, given that we know everything else:

$$F \equiv mx^N + r + Ix^L \pmod{P}$$

- (It turns out $F = 0$ for the adjustable table)

Summary of reverse engineering CRC

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Not guaranteed, but we can find CRC parameters from (m, r) pairs
 - refin and refout, only 4 possibilities
 - width is guessed or based on degree of GCD
 - poly or P determined by taking GCD, and possibly factoring
 - init or I solved with Gaussian elimination
 - May not be unique, can have equivalent (I, F) pairs
 - xorout or F computed from other parameters

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Conclusion

Conclusion

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- CRCs are simple
- Simple stuff can have deep theory behind it
- $\text{GF}(2)$ and $\text{GF}(2)[x]$ sit at the intersection of computer science, ring theory, and linear algebra
 - Useful for checking for transmission/storage errors
 - Useful for cryptography too (AES-GCM)

This talk has been literally everything I know about CRCs...

Things I still don't know

- Types of errors CRC can detect
- Picking CRC polynomials with good error detection properties...

Thanks

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

- Check out my blog post for some other details
- Blog: <https://dram.page/p/crc-tricks>
- Slides: <https://dram.page/p/crc-tricks/crc-tricks.pdf>

References

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Cook, G. (2021a). Catalogue of parametrised CRC algorithms.
<https://reveng.sourceforge.io/crc-catalogue/>.

Cook, G. (2021b). CRC RevEng: arbitrary-precision CRC calculator and algorithm finder. <https://reveng.sourceforge.io>.

Ewing, G. (2010). Reverse-engineering a CRC algorithm. <https://www.cosc.canterbury.ac.nz/greg.ewing/essays/CRC-Reverse-Engineering.html>.

Gopal, V., Ozturk, E., Guilford, J., Wolrich, G., Feghali, W., Dixon, M., and Karakoyunlu, D. (2009). Fast CRC computation for generic polynomials using pclmulqdq instruction.
<https://www.intel.com/content/dam/www/public/us/en/documents/white-papers/fast-crc-computation-generic-polynomials-pclmulqdq-paper.pdf>. Intel White Paper.

References (Cont'd)

CRC tricks

dram

Background

Cyclic
redundancy
check

Efficient CRC
with CLMUL

Preimage
attacks

Reverse
engineering

Conclusion

Kutenin, D. (2021). How a bug(?) in the linux CRC-32 checksum turned out not to be a bug. <https://danlark.org/2021/03/08/how-a-bug-in-the-linux-crc-32-checksum-turned-out-not-to-be-a-bug/>.

Williams, R. N. (1993). A painless guide to CRC error detection algorithms. https://zlib.net/crc_v3.txt.

Wolf, C. (2019). Reference implementations of various CRCs using carry-less multiply. <http://svn.clairexen.net/handicraft/2018/clmulcrc/>.

Wolf, C. (2021). RISC-V Bitmanip extension, document version 0.93. Technical report.